

Here is a short entry for each separate coding project I worked on in ITAI 4370.

1. Digital transmission simulation (Module 1 lab)

- ****Problem statement****

Simulate a simple digital telecom transmission chain to see how bits are generated, modulated, corrupted by noise, and recovered at the receiver.

- ****Methods and tools used****

Python, NumPy, and Matplotlib to generate random bits, map them to symbols, add noise, and compare original vs recovered bit streams.

- ****Code snippet / repo****

```
```python
bits = np.random.randint(0, 2, size=1000)
symbols = 2*bits - 1 # BPSK mapping
noise = np.random.normal(0, 0.5, len(symbols))
rx = symbols + noise
bits_hat = (rx > 0).astype(int)
```
```

2. RF propagation & FSPL for WiFi bands (Lab 2, Module 2)

- ****Problem statement****

Quantify how free space path loss grows with distance for 2.4 GHz vs 5 GHz WiFi to understand wireless coverage tradeoffs.

- ****Methods and tools used****

Python, NumPy, Matplotlib; applied the FSPL formula in dB over a range of distances and plotted both frequency bands.

- ****Code snippet / repo****

```
```python
f1, f2 = 2400, 5000 # MHz
d = np.linspace(0.1, 20, 100) # km
fspl1 = 20*np.log10(d) + 20*np.log10(f1) + 32.44
fspl2 = 20*np.log10(d) + 20*np.log10(f2) + 32.44
plt.plot(d, fspl1, label="2.4 GHz")
plt.plot(d, fspl2, label="5 GHz")
```
```

3. Network traffic prediction (Module 3 lab 1)

- ****Problem statement****

Predict near-term network traffic volume from historical usage patterns to support basic capacity planning.

- **Methods and tools used**

Python, pandas, scikit-learn regression models; engineered features like hour-of-day and day-of-week and evaluated prediction error on a held-out set.

- **Code snippet / repo**

```
```python
X = df[["hour", "dow", "prev_hour"]]
y = df["traffic_mbps"]
model = RandomForestRegressor(n_estimators=200, random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
````
```

4. Network anomaly detection (Module 3 lab 2)

- **Problem statement**

Detect unusual network traffic patterns (potential attacks or faults) from flow-level features.

- **Methods and tools used**

Python, pandas, scikit-learn; trained an unsupervised anomaly detector (e.g., Isolation Forest) on mostly normal traffic and flagged outliers as anomalies.

- **Code snippet / repo**

```
```python
from sklearn.ensemble import IsolationForest

iso = IsolationForest(contamination=0.01, random_state=42)
iso.fit(X_train)
anomaly_scores = iso.decision_function(X_test)
anomalies = iso.predict(X_test) == -1
````
```

5. Customer segmentation (Module 3 lab 3)

- **Problem statement**

Group telecom customers into segments based on usage and behavior for targeted services or marketing.

- **Methods and tools used**

Python, pandas, scikit-learn K-means clustering; normalized numeric features and interpreted cluster centroids to name segments.

- **Code snippet / repo**

```
```python
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans

X_scaled = StandardScaler().fit_transform(X)
kmeans = KMeans(n_clusters=4, random_state=42)
clusters = kmeans.fit_predict(X_scaled)
df["segment"] = clusters
```
---
```

6. Network slicing with demand-based resource allocation (Module 4 lab)

- **Problem statement**

Simulate a simple 5G network slicing scenario where resources are dynamically split between slices based on changing demand.

- **Methods and tools used**

Python and NumPy; modeled total capacity and slice demands, then computed per-slice allocation each time step and tracked utilization.

- **Code snippet / repo**

```
```python
total_bw = 100.0
for t in range(T):
 d = np.array([d_slice1[t], d_slice2[t], d_slice3[t]])
 share = d / (d.sum() + 1e-8)
 alloc = total_bw * share
 usage_log.append(alloc)
```
---
```

7. ORAN traffic prediction with linear regression (Module 5 lab)

- **Problem statement**

Show how a simple ML model can forecast ORAN traffic as a basic building block for AI-assisted RAN optimization.

- **Methods and tools used**

Python, pandas, scikit-learn LinearRegression; used time-based features to fit a quick baseline model.

- **Code snippet / repo**

```
```python
from sklearn.linear_model import LinearRegression

X = df[["hour", "dow"]]
```
---
```

```

y = df["traffic_gbps"]
lr = LinearRegression()
lr.fit(X_train, y_train)
y_pred = lr.predict(X_test)
```

8. Multi-model traffic forecasting: ARIMA, linear regression, LSTM (Lab 4, Module 6)

- **Problem statement**
 Compare classical statistical and ML models for forecasting hourly network traffic on a real-world dataset.

- **Methods and tools used**
 Python, pandas, statsmodels ARIMA, scikit-learn LinearRegression, and Keras LSTM; performed train/test split, feature engineering, and evaluated each model with MAE/MSE/R2.
 (Code in `ITAI_4370_Lab_04_Oren_Moreno.ipynb`.)

- **Code snippet / repo**
```python
# ARIMA fit
arima = ARIMA(train["traffic"], order=(2,1,2)).fit()
arima_forecast = arima.forecast(steps=len(test))

# Simple LSTM skeleton
model = Sequential([
    LSTM(32, input_shape=(seq_len, n_features)),
    Dense(1)
])
model.compile(optimizer="adam", loss="mse")
```

9. Edge AI model compression for IoT sensors (Lab 5, Module 7)

- **Problem statement**
 Take a cloud-scale NN for IoT sensor classification and compress it so it can run efficiently on constrained edge devices.

- **Methods and tools used**
 TensorFlow/Keras for the baseline model, pruning and quantization, and TensorFlow Lite conversion; simulated latency and accuracy for edge vs cloud.
 (Code in `ITAI_4370_Lab_05_Oren_Moreno_2025Fall.ipynb`.)

- **Code snippet / repo**
```python

```

```
base_model = keras.Sequential([
    layers.Dense(64, activation="relu", input_shape=(n_features,)),
    layers.Dense(32, activation="relu"),
    layers.Dense(3, activation="softmax"),
])
```

```
converter = tf.lite.TFLiteConverter.from_keras_model(base_model)
converter.optimizations = [tf.lite.Optimize.DEFAULT]
tflite_model = converter.convert()
```
```

---

### ### 10. ML-based IDS on NSL-KDD (Module 8 lab)

- **\*\*Problem statement\*\***

Build an intrusion detection system that can classify network connections as normal or specific attack types using NSL-KDD.

- **\*\*Methods and tools used\*\***

Python, pandas, scikit-learn RandomForest and SVM classifiers, plus Isolation Forest and One-Class SVM for unsupervised detection; used SHAP for feature importance explanations.

- **\*\*Code snippet / repo\*\***

```
```python
rf = RandomForestClassifier(n_estimators=300, random_state=42)
rf.fit(X_train, y_train)
y_pred = rf.predict(X_test)
```

```
explainer = shap.TreeExplainer(rf)
shap_values = explainer.shap_values(X_test_sample)
```
```

---

### ### 11. Energy forecasting & base-station sleep scheduling (Module 9 lab)

- **\*\*Problem statement\*\***

Forecast telecom energy consumption and design a base-station sleep strategy that reduces cost and emissions while keeping QoS acceptable.

- **\*\*Methods and tools used\*\***

Python, Prophet and LSTM models for energy/time-series prediction; evaluated different sleep windows and computed energy and carbon savings.

- **\*\*Code snippet / repo\*\***

```
```python
# Prophet
m = Prophet()
```

```

m.fit(df_prophet)
forecast = m.predict(future)

# Simple sleep rule
df["sleep"] = (df["pred_traffic"] < threshold).astype(int)
saved_kwh = (df["sleep"] * base_station_power * dt).sum()
```

12. Digital twin of a 3-router network with SimPy (Midterm Q1, Module 10)

- **Problem statement**
 Build a live digital twin of a small network (three routers in series) to study throughput, latency, and queue growth as traffic conditions change.

- **Methods and tools used**
 SimPy for discrete-event simulation, NumPy and statistics for metrics, Matplotlib for time series plots of queue lengths and latency.
 (Code in `ITAI_4370_Midterm_Oren_Moreno_2025Fall.ipynb`.)

- **Code snippet / repo**
```python
def router(env, in_q, out_q, service_time):
    while True:
        pkt = yield in_q.get()
        yield env.timeout(service_time)
        yield out_q.put(pkt)

env = simpy.Environment()
r1_q, r2_q, r3_q = simpy.Store(env), simpy.Store(env), simpy.Store(env)
env.process(router(env, r1_q, r2_q, 0.5))
```

13. Agent-based network modeling with Mesa (Midterm Q2)

- **Problem statement**
 Use an agent-based model to simulate packets flowing through a network graph and observe congestion hotspots and routing behavior.

- **Methods and tools used**
 Mesa for ABM structure, NetworkX for graph topology, and Matplotlib for visualizing node load and packet paths.
 (Code in `ITAI_4370_Midterm_Oren_Moreno_2025Fall.ipynb`.)

- **Code snippet / repo**
```python
class RouterAgent(Agent):

```

```

def __init__(self, unique_id, model):
    super().__init__(unique_id, model)
    self.queue = []

    def step(self):
        if self.queue:
            pkt = self.queue.pop(0)
            next_hop = random.choice(self.model.graph.neighbors(self.unique_id))
            self.model.route_packet(pkt, next_hop)
        ...

```

14. Predictive optimization of routing with ML (Midterm Q3)

- **Problem statement**

Learn a model of future traffic and use it to dynamically split load across two parallel paths in order to minimize total latency compared to a static policy.

- **Methods and tools used**

Python, scikit-learn RandomForestRegressor for single-step traffic prediction; applied a latency cost function and compared predictive vs static routing policies.

- **Code snippet / repo**

```

```python
model = RandomForestRegressor(n_estimators=300, random_state=42)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

s = np.clip(alpha / (y_pred + 1e-3), 0.1, 0.9) # split to path 1
latency_ml = latency(path1, s*y) + latency(path2, (1-s)*y)
```

```

15. Federated learning simulation for 6G resource allocation (Module 11 lab)

- **Problem statement**

Simulate a federated learning setup where multiple edge nodes collaboratively train a model for predictive resource allocation without sharing raw data.

- **Methods and tools used**

Python and PyTorch; each client trained locally on its partitioned data, sent model updates to a central server for weighted averaging (FedAvg), and then evaluated global performance vs a centralized baseline.

- **Code snippet / repo**

```

```python
def client_update(model, loader, epochs=1):
 opt = torch.optim.SGD(model.parameters(), lr=0.01)

```

```
for _ in range(epochs):
 for x, y in loader:
 opt.zero_grad()
 loss = criterion(model(x), y)
 loss.backward()
 opt.step()
return model.state_dict()

Server side aggregation
global_state = average_states(client_states, weights)
```
```